

Neural Networks and Deep Learning-Assignment 5

Name: Sai Swetha Nambari

ID:700742846

Github link: <https://github.com/SwethaNam/nnassign5.git>

Video link:

<https://drive.google.com/file/d/1QJ90Ei2c7kPZXKv5peBz008asLsbEP54/view?usp=sharing>

Question 1:

First we run the code given in the class as shown below before question 1:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import re
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers import Dense, Embedding, LSTM, SpatialDropout1D
from keras.utils.np_utils import to_categorical

[ ] from google.colab import drive
drive.mount('/content/gdrive')
path_to_csv = '/content/gdrive/My Drive/Sentiment.csv'

Mounted at /content/gdrive

[ ] import pandas as pd
dset = pd.read_csv(path_to_csv, header=0)
mask = dset.columns.isin(['text', 'sentiment'])
data = dset.loc[:, mask]
data['text'] = data['text'].apply(lambda x: x.lower())
data['text'] = data['text'].apply(lambda x: re.sub('[^a-zA-z0-9\s]', '', x))
```

Keeping only the necessary columns

In this step we import the libraries and then import the the sentiment.csv from the drive. Later we read the csv file and select only the necessary columns text and sentiment as shown above.

Next we use maximum words on 2000 to tokenize the sentence and then compile the model as shown below:

```

▶ for idx, row in data.iterrows():
    row[0] = row[0].replace('rt', ' ')
max_fatures = 2000
tokenizer = Tokenizer(num_words=max_fatures, split=' ')
tokenizer.fit_on_texts(data['text'].values)
X = tokenizer.texts_to_sequences(data['text'].values)

X = pad_sequences(X) #Padding the feature matrix

embed_dim = 128
lstm_out = 196
def createmodel():
    model = Sequential() #Sequential Neural Network
    model.add(Embedding(max_fatures, embed_dim,input_length = X.shape[1])) #input dimension 2000 Neurons, output dimension 128 Neurons
    model.add(LSTM(lstm_out, dropout=0.2, recurrent_dropout=0.2))
    model.add(Dense(3,activation='softmax'))
    model.compile(loss = 'categorical_crossentropy', optimizer='adam',metrics = ['accuracy']) #Compiling the model
    return model
# print(model.summary())

```

```

[ ] label_encoder = LabelEncoder() #Applying label Encoding on the label matrix
integer_encoded = label_encoder.fit_transform(data['sentiment']) #fitting the model
y = to_categorical(integer_encoded)
X_train, X_test, Y_train, Y_test = train_test_split(X,y, test_size = 0.33, random_state = 42)

batch_size = 32
model = createmodel() #Function call to Sequential Neural Network
model.fit(X_train, Y_train, epochs = 1, batch_size=batch_size, verbose = 2)
score,acc = model.evaluate(X_test,Y_test,verbose=2,batch_size=batch_size) #evaluating the model
print(score)
print(acc)

291/291 - 61s - loss: 0.8245 - accuracy: 0.6424 - 61s/epoch - 210ms/step
144/144 - 3s - loss: 0.7612 - accuracy: 0.6769 - 3s/epoch - 23ms/step
0.7611876726150513
0.6769331693649292

```

```

[ ] print(model.metrics_names)

```

```

['loss', 'accuracy']

```

In the above screenshot, we apply label encoding on the label matrix and then fit the model and lastly evaluate the model.

Then we print the model metrics.

Question 1: Save the model and use the saved model to predict new text data (ex, “A lot of good things are happening. We are respected again throughout the world, and that's a great [thing.@realDonaldTrump](#)”)

```
[ ] #question 1:
    model.save('sentimentAnalysis.h5') #Saving the model

[ ] from keras.models import load_model
    model= load_model('sentimentAnalysis.h5') #loading the saved model
    print(integer_encoded)
    print(data['sentiment'])

[1 2 1 ... 2 0 2]
0      Neutral
1      Positive
2      Neutral
3      Positive
4      Positive
...
13866   Negative
13867   Positive
13868   Positive
13869   Negative
13870   Positive
Name: sentiment, Length: 13871, dtype: object
```

Here we have saved the model, and loaded and printed the saved model.

```
# Predicting on the text data
sentence = ['A lot of good things are happening. We are respected again throughout the world, and that is a great thing.@realDonaldTrump']
sentence = tokenizer.texts_to_sequences(sentence)
sentence = pad_sequences(sentence, maxlen=28, dtype='int32', value=0)
sentiment_probs = model.predict(sentence, batch_size=1, verbose=2)[0]
sentiment = np.argmax(sentiment_probs)

print(sentiment_probs)
if sentiment == 0:
    print("Neutral")
elif sentiment < 0:
    print("Negative")
elif sentiment > 0:
    print("Positive")
else:
    print("Cannot be determined")

1/1 - 0s - 324ms/epoch - 324ms/step
[0.40014437 0.10924453 0.4906111 ]
Positive
```

Predicting on the text data.

Question 2: Apply GridSearchCV on the source code provided in the class

```
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV
model = KerasClassifier(build_fn=create_model, verbose=2)
batch_size = [10, 20, 40]
epochs = [1, 2]
param_grid = {'batch_size': batch_size, 'epochs': epochs} #creating dictionary for batch size, no. of epochs
grid = GridSearchCV(estimator=model, param_grid=param_grid) #Applying dictionary with
grid_result = grid.fit(X_train, Y_train) #Fitting the model

print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_)) #best score, best hyper parameters

372/372 - 62s - loss: 0.6827 - accuracy: 0.7097 - 62s/epoch - 168ms/step
93/93 - 3s - loss: 0.7318 - accuracy: 0.6783 - 3s/epoch - 35ms/step
Epoch 1/2
372/372 - 65s - loss: 0.8365 - accuracy: 0.6403 - 65s/epoch - 175ms/step
Epoch 2/2
372/372 - 62s - loss: 0.6765 - accuracy: 0.7147 - 62s/epoch - 168ms/step
93/93 - 3s - loss: 0.7480 - accuracy: 0.6918 - 3s/epoch - 35ms/step
Epoch 1/2
372/372 - 70s - loss: 0.8342 - accuracy: 0.6437 - 70s/epoch - 187ms/step
Epoch 2/2
372/372 - 68s - loss: 0.6697 - accuracy: 0.7166 - 68s/epoch - 182ms/step
93/93 - 2s - loss: 0.7546 - accuracy: 0.6620 - 2s/epoch - 25ms/step
Epoch 1/2
372/372 - 68s - loss: 0.8312 - accuracy: 0.6443 - 68s/epoch - 184ms/step
Epoch 2/2
372/372 - 64s - loss: 0.6711 - accuracy: 0.7099 - 64s/epoch - 172ms/step
93/93 - 2s - loss: 0.8009 - accuracy: 0.6362 - 2s/epoch - 22ms/step
186/186 - 43s - loss: 0.8446 - accuracy: 0.6381 - 43s/epoch - 233ms/step
17/17 - 1s - loss: 0.7671 - accuracy: 0.6482 - 1s/epoch - 30ms/step
186/186 - 44s - loss: 0.8503 - accuracy: 0.6340 - 44s/epoch - 236ms/step
Epoch 2/2
186/186 - 42s - loss: 0.6878 - accuracy: 0.7057 - 42s/epoch - 224ms/step
47/47 - 2s - loss: 0.7525 - accuracy: 0.6789 - 2s/epoch - 33ms/step
Epoch 1/2
186/186 - 45s - loss: 0.8436 - accuracy: 0.6373 - 45s/epoch - 242ms/step
Epoch 2/2
186/186 - 42s - loss: 0.6903 - accuracy: 0.7044 - 42s/epoch - 226ms/step
47/47 - 3s - loss: 0.7339 - accuracy: 0.6808 - 3s/epoch - 53ms/step
Epoch 1/2
186/186 - 41s - loss: 0.8470 - accuracy: 0.6374 - 41s/epoch - 222ms/step
Epoch 2/2
186/186 - 41s - loss: 0.6861 - accuracy: 0.7053 - 41s/epoch - 220ms/step
47/47 - 1s - loss: 0.7870 - accuracy: 0.6749 - 1s/epoch - 31ms/step
Epoch 1/2
930/930 - 154s - loss: 0.8103 - accuracy: 0.6552 - 154s/epoch - 165ms/step
Epoch 2/2
930/930 - 150s - loss: 0.6683 - accuracy: 0.7156 - 150s/epoch - 162ms/step
Best: 0.679758 using {'batch_size': 10, 'epochs': 2}
```